

## Importing Libraries


```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import scipy
%matplotlib inline
plt.style.use('fivethirtyeight')
pd.set_option('display.max_rows', 500)
pd.set_option('display.max_columns', 500)
pd.set_option('display.width', 1000)
```

## Loading the Dataset :

```
In [2]: data = pd.read_csv('ml_project1_data_pt.csv', delimiter = ',')
data.head(3)
```

Out[2]:

	ID	ano_nasc	educacao	estado_civil	renda_ano	crianca_casa	adoles_casa	dt_primcomp	recencia_dias	vinho_montante	frutas_montante
0	5524	1957	Graduation	Single	58138.0	0	0	9/4/2012	58	635	88
1	2174	1954	Graduation	Single	46344.0	1	1	3/8/2014	38	11	1
2	4141	1965	Graduation	Together	71613.0	0	0	8/21/2013	26	426	49



```
In [3]: data.shape
```

Out[3]: (2240, 27)

```
In [4]: data.columns
```

Out[4]: Index(['ID', 'ano\_nasc', 'educacao', 'estado\_civil', 'renda\_ano', 'crianca\_casa', 'adoles\_casa', 'dt\_primcomp', 'recencia\_dias', 'vinho\_montante', 'frutas\_montante', 'carne\_montante', 'peixe\_montante', 'doces\_montante', 'ouro\_montante', 'promocoes\_desconto', 'promocoes\_web', 'promocoes\_catalogo', 'promocoes\_store', 'num\_visit\_web\_ult\_mes', 'Cmp3', 'Cmp4', 'Cmp5', 'Cmp1', 'Cmp2', 'reclamacoes', 'target'], dtype='object')

# Data Preprocessing

```
In [4]: # Checking for null values.
info = pd.DataFrame(data=data.isnull().sum()).T.rename(index={0: 'Null values'})
info = info.append(pd.DataFrame(data=data.isnull().sum()/data.shape[0] * 100).T.rename(index={0: '% Null values'}))
info
```

Out[4]:

	ID	ano_nasc	educacao	estado_civil	renda_ano	crianca_casa	adoles_casa	dt_primcomp	recencia_dias	vinho_montante	frutas_montant
Null values	0.0	0.0	0.0	0.0	24.000000	0.0	0.0	0.0	0.0	0.0	0.
% Null values	0.0	0.0	0.0	0.0	1.071429	0.0	0.0	0.0	0.0	0.0	0.

```
In [5]: # Checking for Duplicates :
data.duplicated().sum()
```

Out[5]: 0

```
In [6]: data.describe()
```

Out[6]:

	ID	ano_nasc	renda_ano	crianca_casa	adoles_casa	recencia_dias	vinho_montante	frutas_montante	carne_montante	peix
count	2240.000000	2240.000000	2216.000000	2240.000000	2240.000000	2240.000000	2240.000000	2240.000000	2240.000000	
mean	5592.159821	1968.805804	52247.251354	0.444196	0.506250	49.109375	303.935714	26.302232	166.950000	
std	3246.662198	11.984069	25173.076661	0.538398	0.544538	28.962453	336.597393	39.773434	225.715373	
min	0.000000	1893.000000	1730.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	
25%	2828.250000	1959.000000	35303.000000	0.000000	0.000000	24.000000	23.750000	1.000000	16.000000	
50%	5458.500000	1970.000000	51381.500000	0.000000	0.000000	49.000000	173.500000	8.000000	67.000000	
75%	8427.750000	1977.000000	68522.000000	1.000000	1.000000	74.000000	504.250000	33.000000	232.000000	
max	11191.000000	1996.000000	666666.000000	2.000000	2.000000	99.000000	1493.000000	199.000000	1725.000000	

```
In [7]: data['dt_primcomp'] = pd.to_datetime(data['dt_primcomp'], errors='coerce')
data['dt_primcomp'] = data['dt_primcomp'].dt.strftime('%m/%Y')
```

```
In [8]: data['age']= 2020 - data['ano_nasc']
```

```
In [9]: data['renda_mes_media']= data['renda_ano']/12
```

```
In [10]: data['campaing_engagement'] =( data['Cmp1']+data['Cmp2']+data['Cmp3']+data['Cmp4']+data['Cmp5']) /5
```

```
In [11]: data['target'] = data['target'].astype(str)
data['Cmp1'] = data['Cmp1'].astype(str)
data['Cmp2'] = data['Cmp2'].astype(str)
data['Cmp3'] = data['Cmp3'].astype(str)
data['Cmp4'] = data['Cmp4'].astype(str)
data['Cmp5'] = data['Cmp5'].astype(str)
data['reclamacoes'] = data['reclamacoes'].astype(str)
data['digital_profile'] = '0'
data['digital_profile'][(data['num_visit_web_ult_mes']< 5) & (data['promocoes_web']<3)]= '1'
```

## Exploratory Data Analysis :

```
In [13]: data.dtypes.groupby(data.dtypes).size()
```

```
Out[13]: int64      17
float64      3
object       11
dtype: int64
```

```
In [15]: dtypes = pd.DataFrame(data.dtypes.rename('type')).reset_index().astype('str')
dtypes = dtypes.query('index != "dt_primcomp",')
dtypes = dtypes.query('index != "ID"')
dtypes = dtypes.query('index != "target"')
numeric = dtypes[(dtypes.type.isin(['int64', 'float64']))]['index'].values
categorical = dtypes[~(dtypes['index'].isin(numeric)) & (dtypes['index'] != 'target']]['index'].values

print('Numeric:\n', numeric)
print('Categorical:\n', categorical)
```

Numeric:

```
['ano_nasc' 'renda_ano' 'crianca_casa' 'adoles_casa' 'recencia_dias'
'vinho_montante' 'frutas_montante' 'carne_montante' 'peixe_montante'
'doces_montante' 'ouro_montante' 'promocoes_desconto' 'promocoes_web'
'promocoes_catalogo' 'promocoes_store' 'num_visit_web_ult_mes' 'age'
'renda_mes_media' 'campaign_engagement']
```


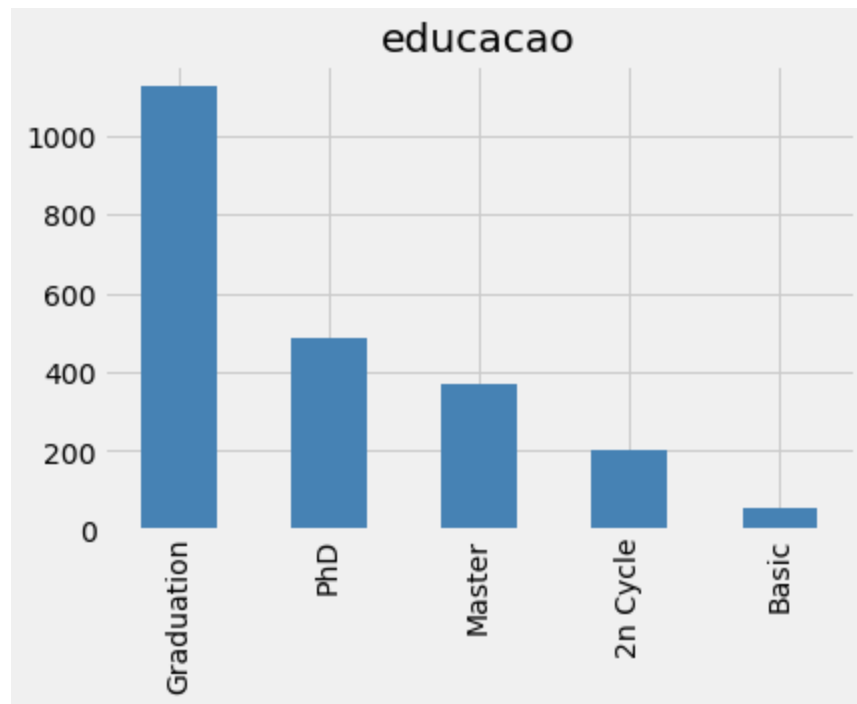
Categorical:



```
['educacao' 'estado_civil' 'Cmp3' 'Cmp4' 'Cmp5' 'Cmp1' 'Cmp2'
'reclamacoes' 'digital_profile']
```



## Categorical Data Analysis

```
In [830]: pylab.rcParams['figure.figsize'] = (6.0, 4.0)
```

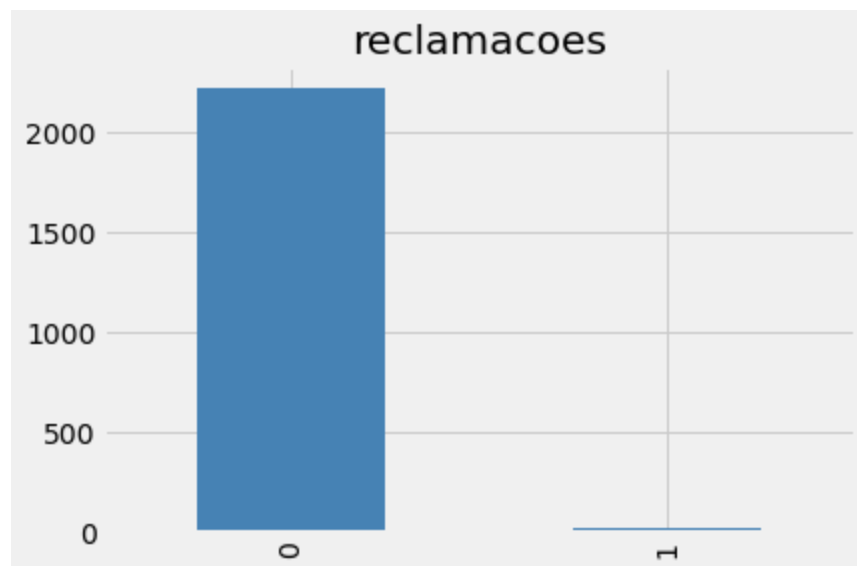

```
In [16]: for attr in categorical:
         figsize=(8,4)
         plt.figure()
         data[attr].value_counts().plot(kind='bar', color='steelblue');
         plt.title(attr);
```














```
In [56]: for attr in categorical:
    fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(12,4))
    outcome_counts = data.groupby([attr, 'target']).size().rename('count').reset_index()
    by_outcome = outcome_counts.pivot(columns='target', values='count', index=attr)
    # Plot the proportions
    by_outcome.div(by_outcome.sum(1), axis=0).plot.bar(stacked=True, ax=ax1);
    # Plot the counts
    data[attr].value_counts().plot.bar(ax=ax2, legend=False,color='steelblue');
    print('Support (%s)\n' % attr)
    print(data[attr].value_counts(), '\n')
    plt.title(attr);
```

Support (educacao)

Graduation 1127

PhD 486

Master 370

2n Cycle 203

Basic 54

Name: educacao, dtype: int64

Support (estado\_civil)

Married 864

Together 580

Single 480

Divorced 232

Widow 77

Alone 3

YOLO 2

Absurd 2

Name: estado\_civil, dtype: int64

Support (Cmp3)

0 2077

1 163

Name: Cmp3, dtype: int64

Support (Cmp4)

0 2073

1 167

Name: Cmp4, dtype: int64

Support (Cmp5)

0 2077

1 163

Name: Cmp5, dtype: int64

Support (Cmp1)

0 2096

1 144

Name: Cmp1, dtype: int64

Support (Cmp2)

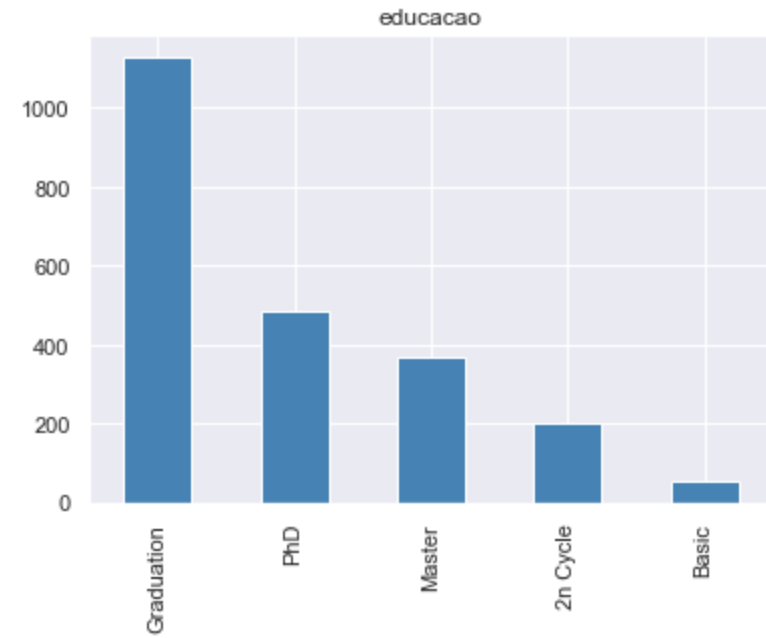

```
0    2210
1       30
Name: Cmp2, dtype: int64
```



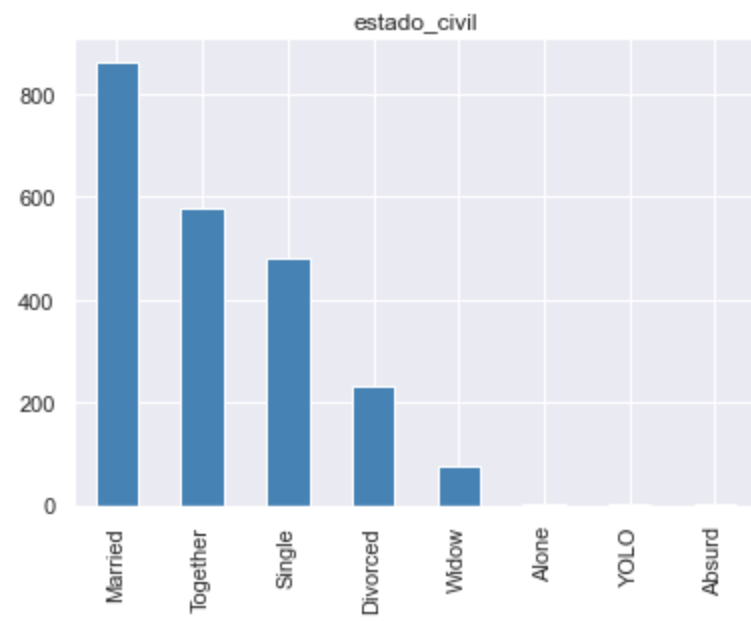

Support (reclamacoes)


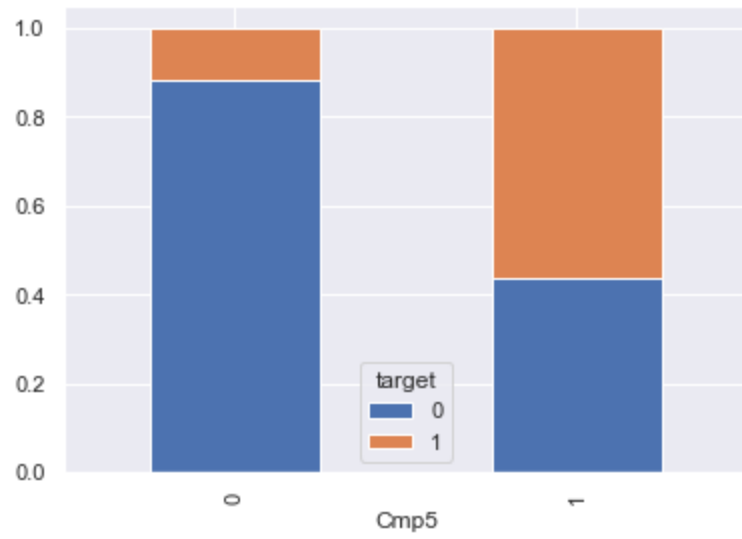


```
0    2219
1       21
Name: reclamacoes, dtype: int64
```





Support (digital\_profile)

```
0    2071
1     169
Name: digital_profile, dtype: int64
```



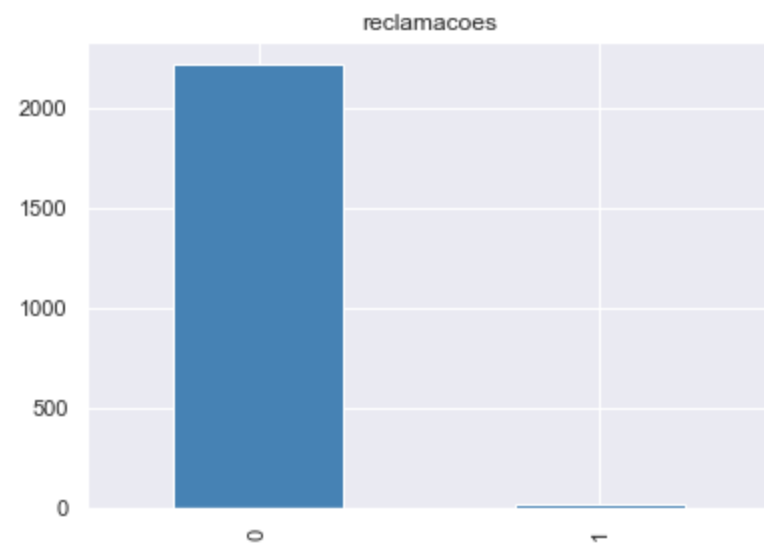
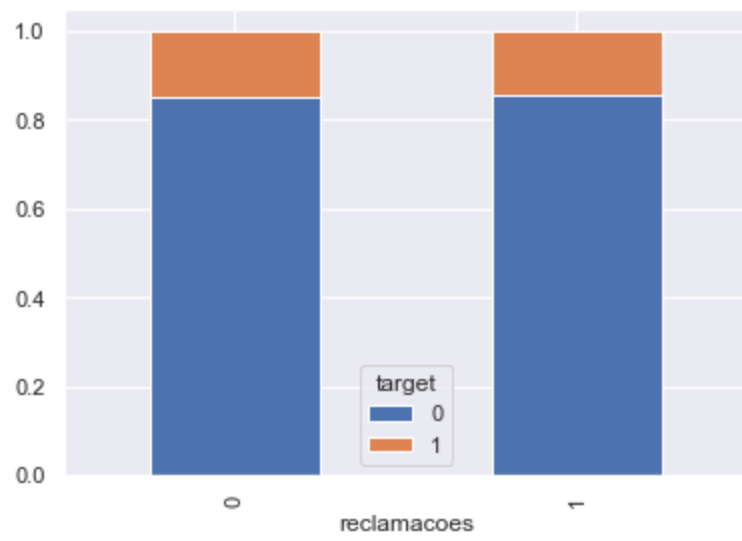














Date Analysis

```
In [18]: data['dt_primcomp'] = pd.to_datetime(data['dt_primcomp'], infer_datetime_format=True)
data.groupby('dt_primcomp')['ID'].nunique().plot(kind='bar')
plt.show()
```




```
In [21]: data['target'] = data['target'].astype(int)
data.groupby('dt_primcomp')['target'].sum().plot(kind='bar')
plt.show()
```



```
In [280]: # campaing
pylab.rcParams['figure.figsize'] = (28, 3)
data.groupby(('dt_primcomp'))['Cmp1', 'Cmp2', 'Cmp3', 'Cmp4', 'Cmp5'].sum().plot(kind='bar')
plt.title("Campaing Success")
plt.figure( figsize=(20, 18))
plt.show()
```

C:\Users\patri\anaconda3\lib\site-packages\ipykernel\_launcher.py:3: FutureWarning: Indexing with multiple keys (implicitly converted to a tuple of keys) will be deprecated, use a list instead.

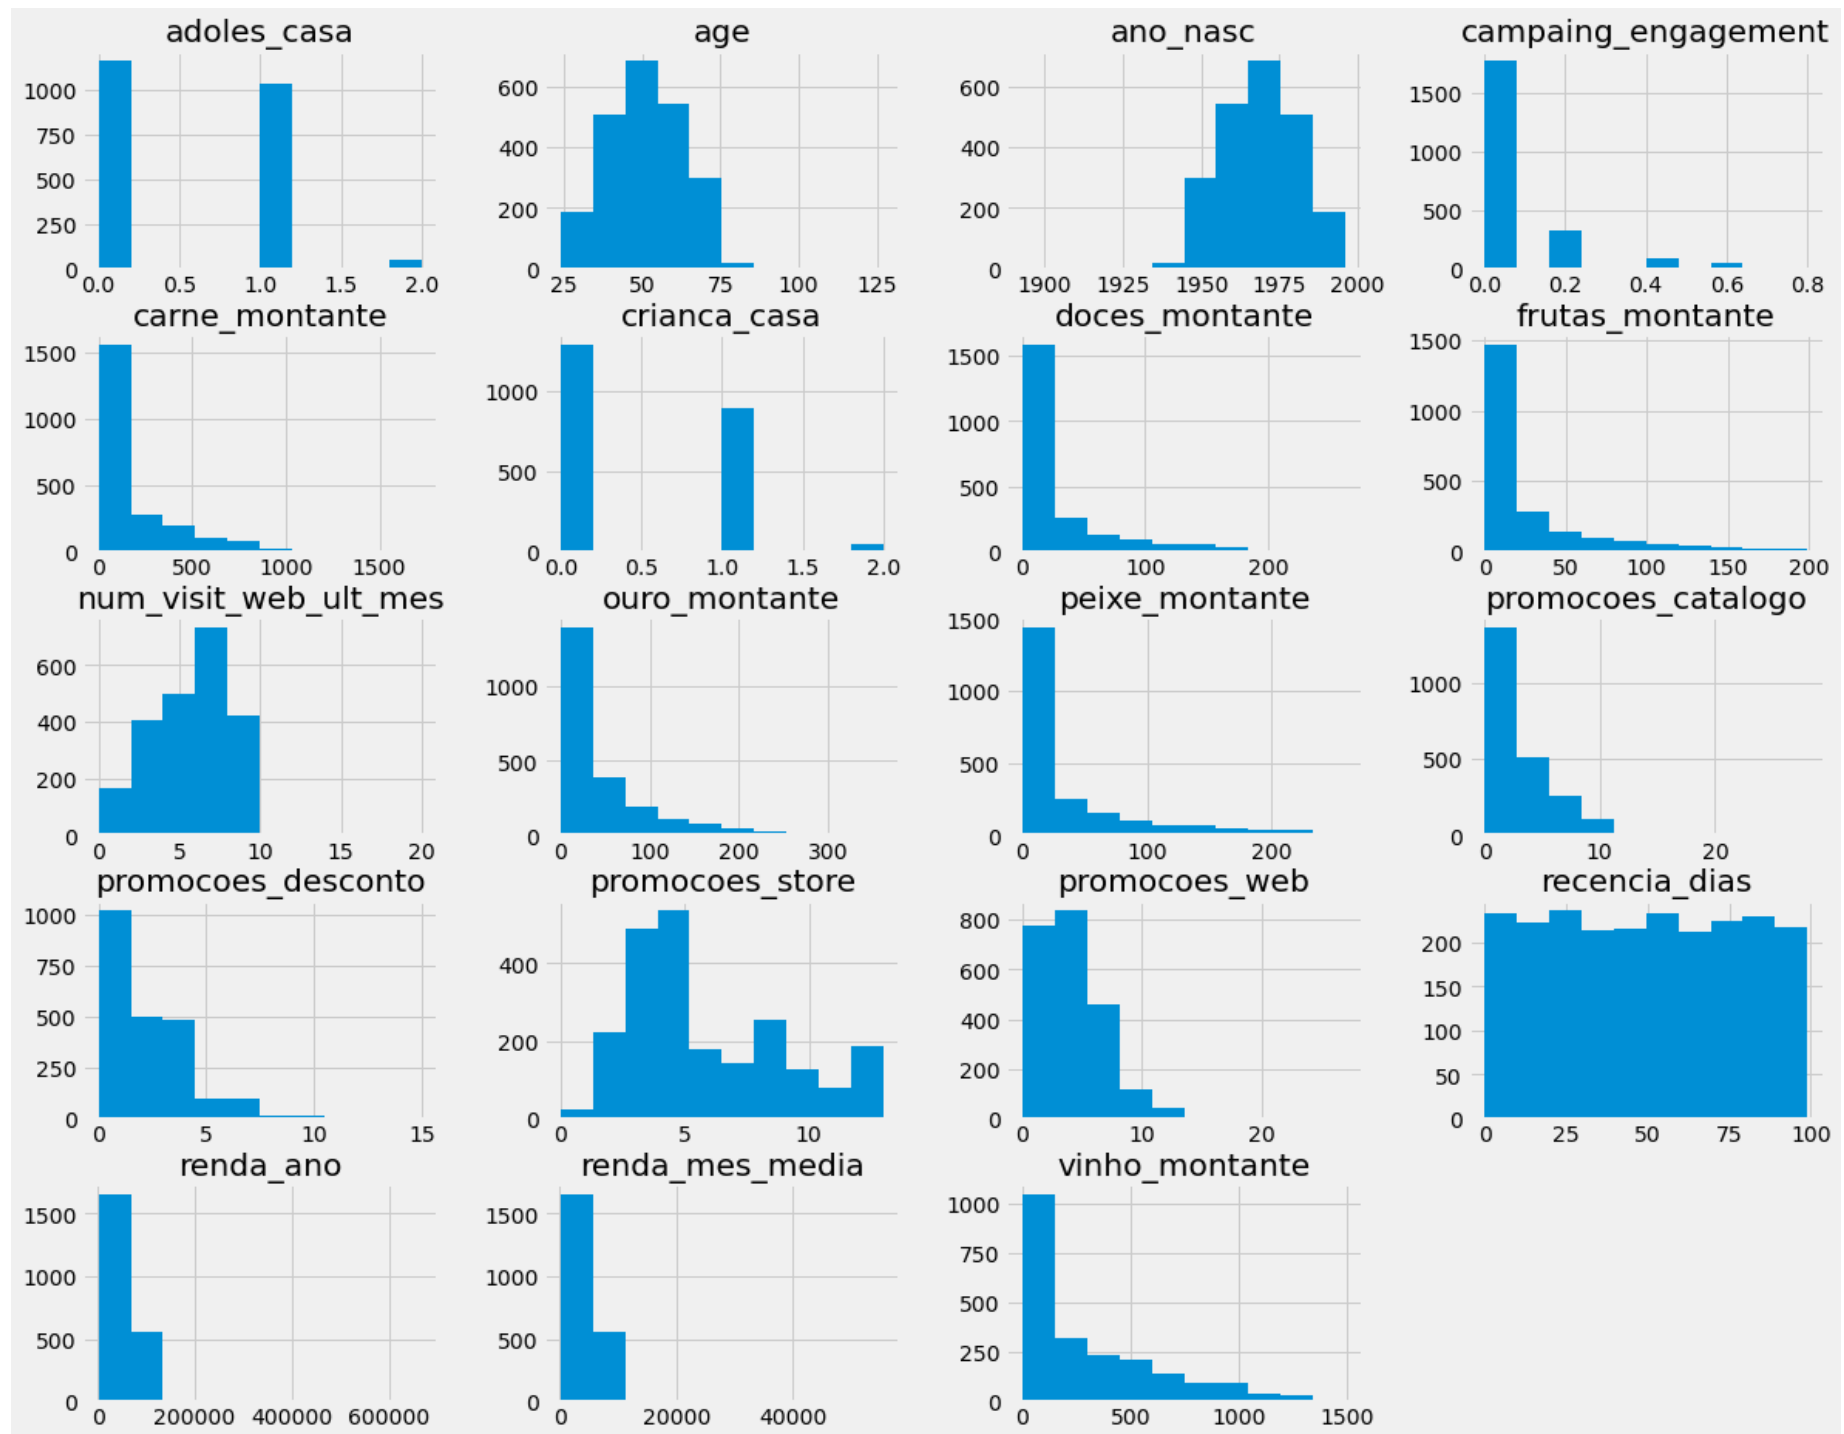
This is separate from the ipykernel package so we can avoid doing imports until



<Figure size 1440x1296 with 0 Axes>

## Numerical Data Analysis

```
In [22]: data[numeric].hist(figsize=(18,15));
```

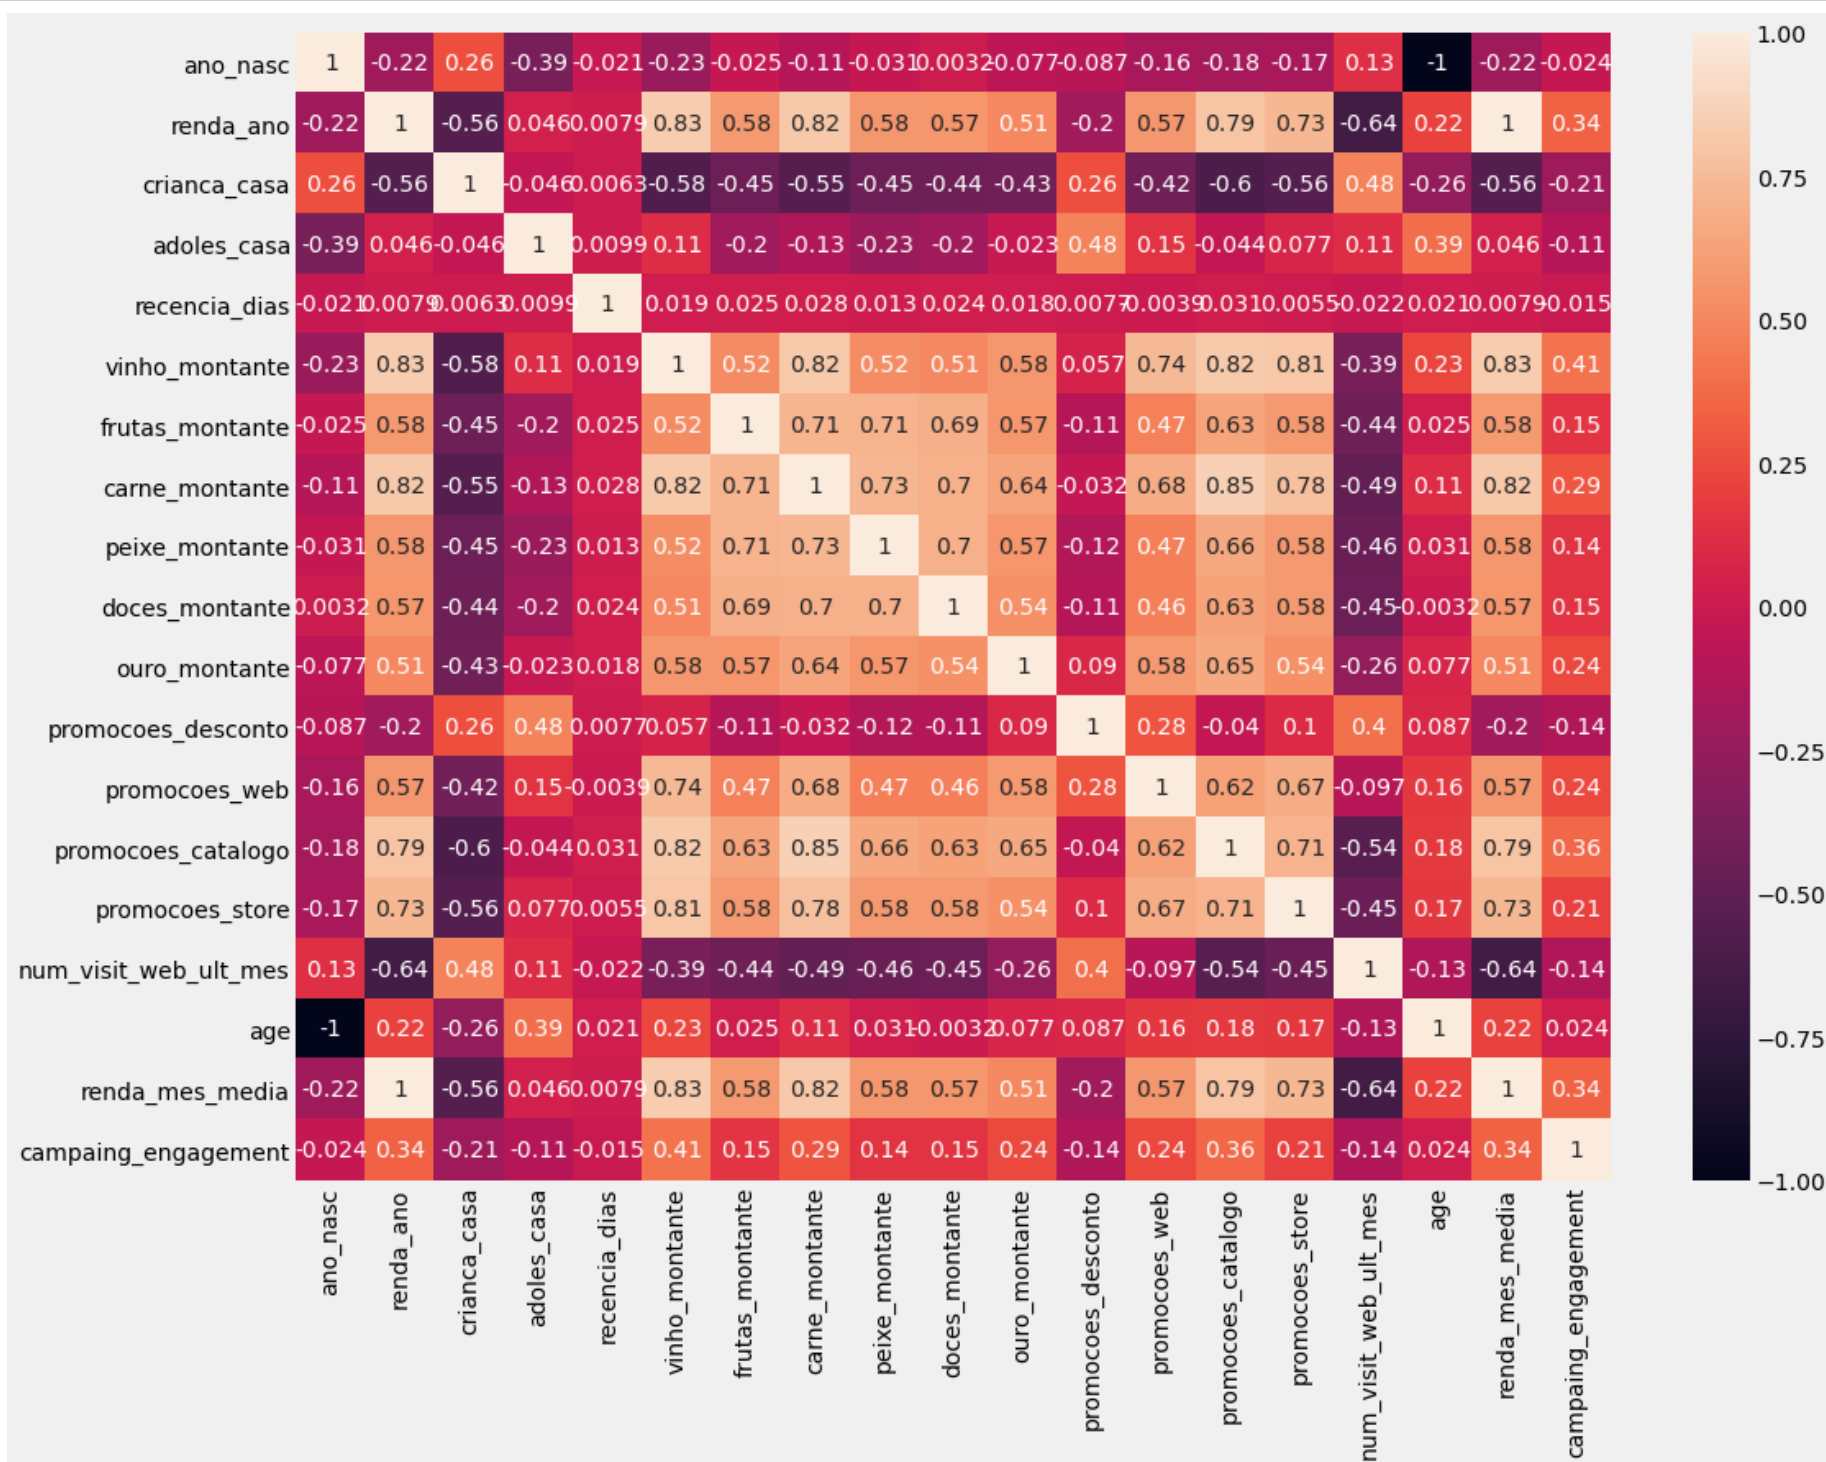


In [23]: data[numeric].describe()

Out[23]:

	ano_nasc	renda_ano	crianca_casa	adoles_casa	recencia_dias	vinho_montante	frutas_montante	carne_montante	peixe_montante	c
count	2240.000000	2216.000000	2240.000000	2240.000000	2240.000000	2240.000000	2240.000000	2240.000000	2240.000000	
mean	1968.805804	52247.251354	0.444196	0.506250	49.109375	303.935714	26.302232	166.950000	37.525446	
std	11.984069	25173.076661	0.538398	0.544538	28.962453	336.597393	39.773434	225.715373	54.628979	
min	1893.000000	1730.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	
25%	1959.000000	35303.000000	0.000000	0.000000	24.000000	23.750000	1.000000	16.000000	3.000000	
50%	1970.000000	51381.500000	0.000000	0.000000	49.000000	173.500000	8.000000	67.000000	12.000000	
75%	1977.000000	68522.000000	1.000000	1.000000	74.000000	504.250000	33.000000	232.000000	50.000000	
max	1996.000000	666666.000000	2.000000	2.000000	99.000000	1493.000000	199.000000	1725.000000	259.000000	

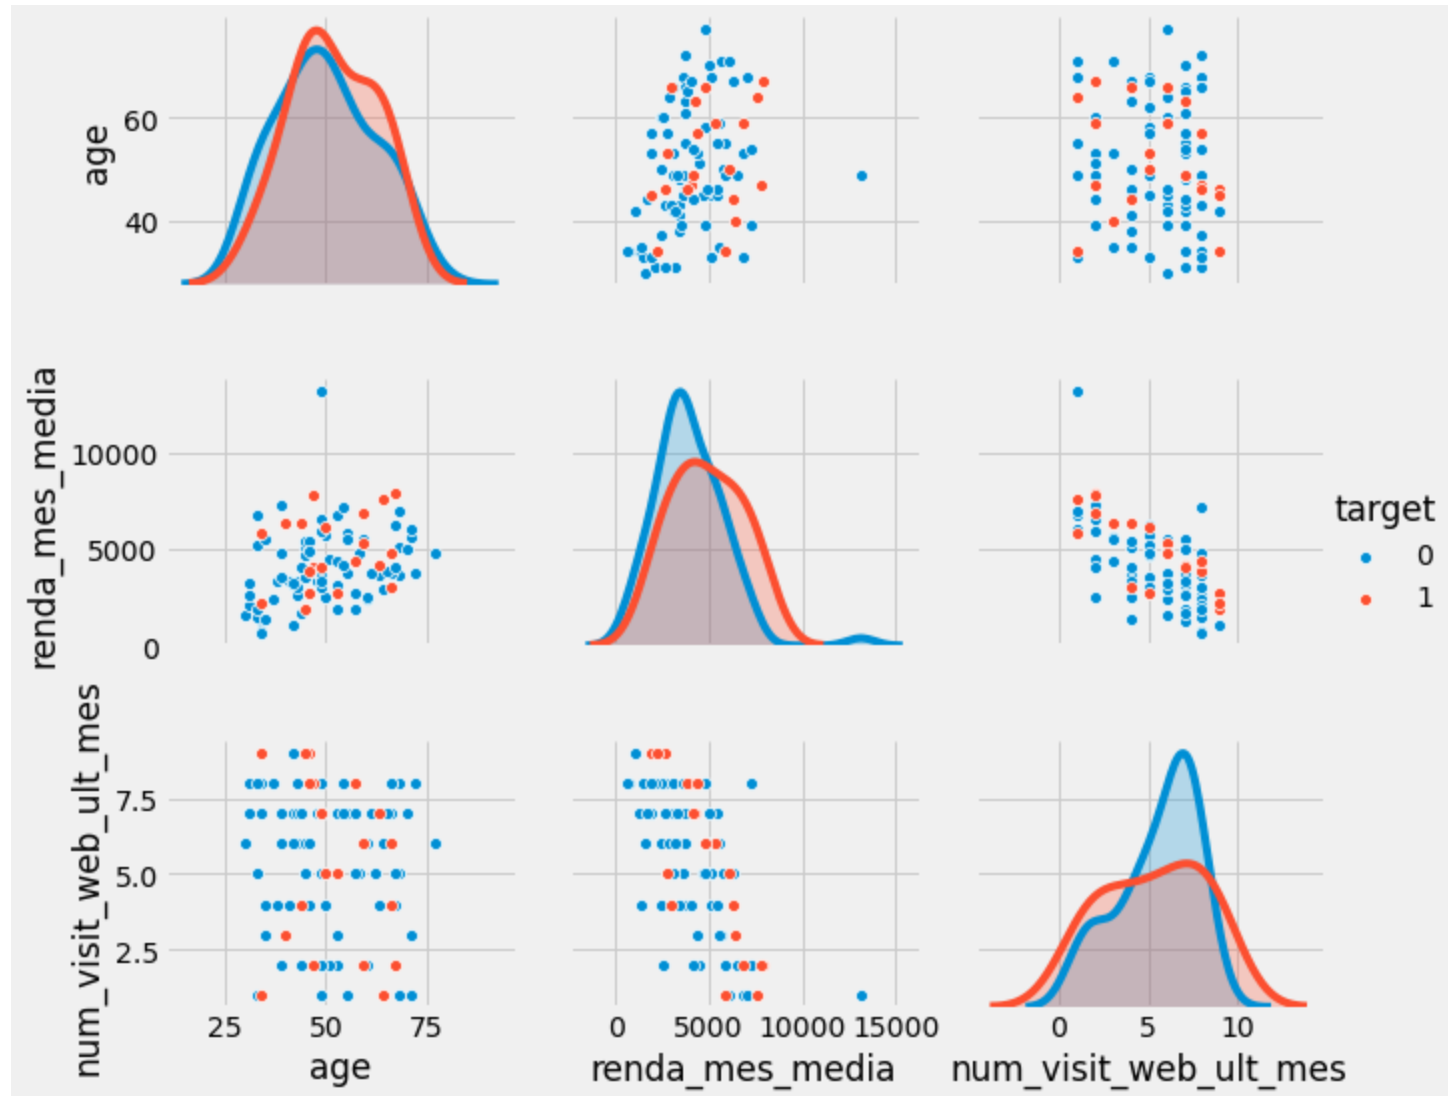
```
In [24]: plt.figure(figsize=(16,12));
sns.heatmap(data[numeric].corr('spearman'), annot=True);
```



## Customer Attributes

```
In [37]: cust_attrs = ['age', 'renda_mes_media', 'num_visit_web_ult_mes', 'target']
```

```
In [38]: data['target'] = data['target'].astype(str)
numeric_outcome = pd.concat([data[numeric], data['target']], axis=1)
sns.pairplot(numeric_outcome[cust_attrs].sample(n=100), hue='target', aspect=1.2);
```



## Clustering



## 1 . Kmeans

```
In [14]: from sklearn.cluster import KMeans
from sklearn.metrics import silhouette_score
from sklearn.preprocessing import StandardScaler
```

```
In [16]: X=data[['renda_ano', 'crianca_casa', 'adoles_casa', 'recencia_dias', 'vinho_montante', 'frutas_montante', 'carne_monta
nte', 'peixe_montante', 'doces_montante', 'ouro_montante', 'promocoes_desconto', 'promocoes_web', 'promocoes_catalogo',
'promocoes_store', 'num_visit_web_ult_mes','age', 'renda_mes_media','campaign_engagement']]
```

```
In [17]: X=X.fillna(0)
```

```
In [48]: # define standard scaler
scaler = StandardScaler()
# transform data
scaled = scaler.fit_transform(X)
print(scaled)
```

```
[[ 0.25193856 -0.82521765 -0.92989438 ...  0.98534473  0.25193856
 -0.43903713]
 [-0.20869932  1.03255877  0.90693402 ...  1.23573295 -0.20869932
 -0.43903713]
 [ 0.77823121 -0.82521765 -0.92989438 ...  0.3176428  0.77823121
 -0.43903713]
 ...
 [ 0.20674965 -0.82521765 -0.92989438 ... -1.01776106  0.20674965
  1.03539042]
 [ 0.68574431 -0.82521765  0.90693402 ...  1.06880747  0.68574431
 -0.43903713]
 [ 0.04614739  1.03255877  0.90693402 ...  1.23573295  0.04614739
 -0.43903713]]
```

```
In [49]: for n_clusters in range(3, 10):
          kmeans = KMeans(init='k-means++', n_clusters = n_clusters, n_init = 30)
          kmeans.fit(scaled)
          clusters = kmeans.predict(scaled)
          sil_avg = silhouette_score(scaled, clusters)
          print("For n_clusters : ", n_clusters, "The average silhouette_score is : ", sil_avg)
```

```
For n_clusters : 3 The average silhouette_score is : 0.2261032879931275
For n_clusters : 4 The average silhouette_score is : 0.15081907973264772
For n_clusters : 5 The average silhouette_score is : 0.15246532108308206
For n_clusters : 6 The average silhouette_score is : 0.15273556967922475
For n_clusters : 7 The average silhouette_score is : 0.15382971566043036
For n_clusters : 8 The average silhouette_score is : 0.15351880625259764
For n_clusters : 9 The average silhouette_score is : 0.14400645664857872
```

```
In [54]: # Choosing number of clusters as 3:
          # Trying Improving the silhouette_score :
          n_clusters = 3
          sil_avg = -1
          while sil_avg < 0.145:
              kmeans = KMeans(init = 'k-means++', n_clusters = n_clusters, n_init = 30)
              kmeans.fit(scaled)
              clusters = kmeans.predict(scaled)
              sil_avg = silhouette_score(scaled, clusters)
              print("For n_clusters : ", n_clusters, "The average silhouette_score is : ", sil_avg)
```

```
For n_clusters : 3 The average silhouette_score is : 0.22575793720918896
```

```
In [55]: # Printing number of elements in each cluster :
          pd.Series(clusters).value_counts()
```

```
Out[55]: 2    1008
          1     644
          0     588
          dtype: int64
```

## Analysing 3 Cluster

```
In [56]: def graph_component_silhouette(n_clusters, lim_x, mat_size, sample_silhouette_values, clusters):
import matplotlib as mpl
mpl.rc('patch', edgecolor = 'dimgray', linewidth = 1)

fig, ax1 = plt.subplots(1, 1)
fig.set_size_inches(8, 8)
ax1.set_xlim([lim_x[0], lim_x[1]])
ax1.set_ylim([0, mat_size + (n_clusters + 1) * 10])
y_lower = 10


for i in range(n_clusters):
    ith_cluster_silhouette_values = sample_silhouette_values[clusters == i]
    ith_cluster_silhouette_values.sort()
    size_cluster_i = ith_cluster_silhouette_values.shape[0]
    y_upper = y_lower + size_cluster_i

    ax1.fill_betweenx(np.arange(y_lower, y_upper), 0, ith_cluster_silhouette_values, alpha = 0.8)

    ax1.text(-0.03, y_lower + 0.5 * size_cluster_i, str(i), color = 'red', fontweight = 'bold',
            bbox = dict(facecolor = 'white', edgecolor = 'black', boxstyle = 'round, pad = 0.3'))

    y_lower = y_upper + 10
```

```
In [57]: # Plotting the intra cluster silhouette distances.  
from sklearn.metrics import silhouette_samples  
sample_silhouette_values = silhouette_samples(scaled, clusters)  
graph_component_silhouette(n_clusters, [-0.07, 0.33], len(X), sample_silhouette_values, clusters)
```



## Dimensionality Analysis


PCA

```
In [58]: from sklearn.decomposition import PCA
```

```
In [59]: pca = PCA()
pca.fit(scaled)
pca_samples = pca.transform(scaled)
```

```
In [60]: # Checking the amount of variance explained :
fig, ax = plt.subplots(figsize=(14, 5))
sns.set(font_scale=1)
plt.step(range(scaled.shape[1]), pca.explained_variance_ratio_.cumsum(), where = 'mid', label = 'Cumulative Variance Explained')
sns.barplot(np.arange(1, scaled.shape[1] + 1), pca.explained_variance_ratio_, alpha = 0.5, color = 'g',
            label = 'Individual Variance Explained')
plt.xlim(0, 20)
plt.xticks(rotation = 45, fontsize = 16)
ax.set_xticklabels([s for s in ax.get_xticklabels()])

plt.ylabel("Explained Variance", fontsize = 18)
plt.xlabel("Principal Components", fontsize = 18)
plt.legend(loc = 'upper right', fontsize = 12)
plt.show()
```




```
In [61]: data['fit_segmentacao'] = kmeans.labels_
```

```
In [62]: data
```

```
Out[62]:
```

	ID	ano_nasc	educacao	estado_civil	renda_ano	crianca_casa	adoles_casa	dt_primcomp	recencia_dias	vinho_montante	frutas_monta
0	5524	1957	Graduation	Single	58138.0	0	0	09/2012	58	635	
1	2174	1954	Graduation	Single	46344.0	1	1	03/2014	38	11	
2	4141	1965	Graduation	Together	71613.0	0	0	08/2013	26	426	
3	6182	1984	Graduation	Together	26646.0	1	0	02/2014	26	11	
4	5324	1981	PhD	Married	58293.0	1	0	01/2014	94	173	
...	...	...	...	...	...	...	...	...	...	...	...
2235	10870	1967	Graduation	Married	61223.0	0	1	06/2013	46	709	
2236	4001	1946	PhD	Together	64014.0	2	1	06/2014	56	406	
2237	7270	1981	Graduation	Divorced	56981.0	0	0	01/2014	91	908	
2238	8235	1956	Master	Together	69245.0	0	1	01/2014	8	428	
2239	9405	1954	PhD	Married	52869.0	1	1	10/2012	40	84	

2240 rows × 32 columns



```
In [63]: data['fit_segmentacao'] = data['fit_segmentacao'].astype(str)
```

## Fit Segmentação Analysis

```
In [64]: from pandas_profiling import ProfileReport
profile = ProfileReport(data, title="Data Profile Report")
```

In [65]: profile





# Overview

Dataset statistics		Variable types	
Number of variables	32	NUM	18
Number of observations	2240	CAT	14
Missing cells	48		
Missing cells (%)	0.1%		
Duplicate rows	0		
Duplicate rows (%)	0.0%		
Total size in memory	560.1 KiB		
Average record size in memory	256.1 B		


## Reproduction

Analysis started	2020-07-21 00:54:36.254745
Analysis finished	2020-07-21 00:56:00.408575
Duration	1 minute and 24.15 seconds
Version	pandas-profiling v2.8.0 ( <a href="https://github.com/pandas-profiling/pandas-profiling">https://github.com/pandas-profiling/pandas-profiling</a> )
Command line	<code>pandas_profiling --config_file config.yaml [YOUR_FILE.csv]</code>

Out[65]:

```
In [188]: # Customer
data.groupby(['fit_segmentacao']).fit_segmentacao.count().sort_values().plot(kind='bar')
data.groupby(['fit_segmentacao']).fit_segmentacao.count().sort_values()
```




```
Out[188]: fit_segmentacao
0         588
1         644
2        1008
Name: fit_segmentacao, dtype: int64
```






```
In [176]: num_cluster = data[['renda_mes_media', 'age', 'crianca_casa', 'adoles_casa', 'recencia_dias', 'num_visit_web_ult_mes', 'campaing_engagement']]
```

```
In [177]: for att in num_cluster:
          figsize=(8,4)
          plt.figure()
          data[att].hist(by=data[ 'fit_segmentacao' ],figsize=(12,9))
          plt.title(att);
```




<Figure size 576x396 with 0 Axes>






<Figure size 576x396 with 0 Axes>






<Figure size 576x396 with 0 Axes>



<Figure size 576x396 with 0 Axes>






<Figure size 576x396 with 0 Axes>






<Figure size 576x396 with 0 Axes>





<Figure size 576x396 with 0 Axes>



```
In [178]: for att in num_cluster:
           print (att)

           print( data.groupby('fit_segmentacao', as_index=True)[att].describe())
```

renda\_mes\_media

	count	mean	std	min	25%	50%	75%	max
fit_segmentacao								
0	586.0	6465.720563	2308.773067	203.916667	5836.500000	6369.041667	6860.8750	55555.500000
1	638.0	4823.648119	946.464420	369.000000	4230.937500	4837.958333	5417.5625	13533.083333
2	992.0	2804.360887	964.706871	144.166667	2116.416667	2836.208333	3470.6250	5541.916667

age

	count	mean	std	min	25%	50%	75%	max
fit_segmentacao								
0	588.0	51.661565	13.646152	25.0	41.0	51.0	63.0	121.0
1	644.0	55.582298	10.127234	28.0	47.0	55.0	64.0	127.0
2	1008.0	48.118056	11.107355	24.0	40.0	47.0	55.0	120.0

crianca\_casa

	count	mean	std	min	25%	50%	75%	max
fit_segmentacao								
0	588.0	0.044218	0.205753	0.0	0.0	0.0	0.0	1.0
1	644.0	0.240683	0.452562	0.0	0.0	0.0	0.0	2.0
2	1008.0	0.807540	0.486830	0.0	1.0	1.0	1.0	2.0

adoles\_casa

	count	mean	std	min	25%	50%	75%	max
fit_segmentacao								
0	588.0	0.171769	0.386421	0.0	0.0	0.0	0.0	2.0
1	644.0	0.936335	0.421782	0.0	1.0	1.0	1.0	2.0
2	1008.0	0.426587	0.518350	0.0	0.0	0.0	1.0	2.0

recencia\_dias

	count	mean	std	min	25%	50%	75%	max
fit_segmentacao								
0	588.0	49.642857	29.420786	0.0	24.00	52.0	74.25	99.0
1	644.0	48.785714	28.658998	0.0	24.75	50.0	72.00	99.0
2	1008.0	49.004960	28.910418	0.0	24.00	49.0	75.00	99.0

num\_visit\_web\_ult\_mes


	count	mean	std	min	25%	50%	75%	max
fit_segmentacao								
0	588.0	2.835034	1.791043	0.0	1.0	2.0	4.0	9.0
1	644.0	5.680124	1.873157	0.0	4.0	6.0	7.0	9.0
2	1008.0	6.531746	1.955560	0.0	5.0	7.0	8.0	20.0

campaing\_engagement

	count	mean	std	min	25%	50%	75%	max
fit_segmentacao								
0	588.0	0.142517	0.204286	0.0	0.0	0.0	0.2	0.8
1	644.0	0.049379	0.106913	0.0	0.0	0.0	0.0	0.6
2	1008.0	0.017659	0.058834	0.0	0.0	0.0	0.0	0.4

```
In [187]: data['digital_profile'] = data['digital_profile'].astype(int)
data.groupby(['fit_segmentacao'])['digital_profile'].sum().plot.bar()
data.groupby(['fit_segmentacao'])['digital_profile'].sum()
```

```
Out[187]: fit_segmentacao
0        37
1        14
2       118
Name: digital_profile, dtype: int32
```



**Classifying the Customers :**

```
In [198]: from sklearn.model_selection import GridSearchCV
from sklearn.metrics import accuracy_score
class Class_Fit(object):
    def __init__(self, clf, params = None):
        if params:
            self.clf = clf(**params)
        else:
            self.clf = clf()

    def train(self, x_train, y_train):
        self.clf.fit(x_train, y_train)

    def predict(self, x):
        return self.clf.predict(x)

    def grid_search(self, parameters, Kfold):
        self.grid = GridSearchCV(estimator = self.clf, param_grid = parameters, cv = Kfold)

    def grid_fit(self, X, Y):
        self.grid.fit(X, Y)

    def grid_predict(self, X, Y):
        self.predictions = self.grid.predict(X)
        print("Precision: {:.2f} %".format(100 * accuracy_score(Y, self.predictions)))
```

```
In [197]: data['target'].dtypes
```

```
Out[197]: dtype('O')
```

```
In [206]: data['target'] = data['target'].astype(str)
```

```
In [210]: columns = ['renda_mes_mediana', 'age', 'recencia_dias', 'vinho_montante', 'frutas_montante', 'carne_montante', 'peixe_montante', 'doces_montante', 'ouro_montante', 'promocoes_desconto', 'promocoes_web', 'promocoes_catalogo', 'promocoes_store', 'num_visitas_web_ult_mes', 'Cmp3', 'Cmp4', 'Cmp5', 'Cmp1', 'Cmp2', 'reclamacoes', 'fit_segmentacao', 'campaign_engagement']
X = data[columns]
Y = data['target']
```

```
In [72]: Y=Y.astype(int)
n_instances = len(X)
p_instances = Y.sum() / len(Y)
p_targeted = 0.15
n_targeted = int(n_instances*p_targeted)

print('Number of instances: {:,}'.format(n_instances))
print('Number of conversions {:,}'.format(Y.sum()))
print('Conversion rate: {:.2f}%'.format(p_instances*100.))
print('15% of the population {:,}'.format(n_targeted))
print('Expected number of conversions targetting {:,} @ {:.2f}%: {:,}'.format(n_targeted, p_instances*100., int(p_instances * n_targeted)))
```

Number of instances: 2,240  
Number of conversions 334  
Conversion rate: 14.91%  
15% of the population 336  
Expected number of conversions targetting 336 @ 14.91%: 50

## Train, Test Splitting

```
In [208]: from sklearn.model_selection import train_test_split
```

```
In [213]: X=X.fillna(0)
X=X.astype(int)
Y=Y.astype(str)
```

```
In [214]: X_train, X_test, Y_train, Y_test = train_test_split(X, Y, train_size = 0.8)
```

## Training Models :

```
In [215]: from sklearn.svm import LinearSVC
from warnings import simplefilter
from sklearn.exceptions import ConvergenceWarning
simplefilter("ignore", category=ConvergenceWarning)
```

```
In [216]: svc = Class_Fit(clf=LinearSVC)
svc.grid_search(parameters = [{ 'C':np.logspace(-2,2,10)}], Kfold = 5)
```

```
In [217]: svc.grid_fit(X=X_train, Y=Y_train)
```

```
In [218]: svc.grid_predict(X_test, Y_test)
```

Precision: 85.04 %

```
In [219]: from sklearn.metrics import confusion_matrix
```

```
In [220]: class_names = [i for i in range(1,11)]  
cnf = confusion_matrix(Y_test, svc.predictions)  
cnf
```

```
Out[220]: array([[380,  1],  
                [ 66,  1]], dtype=int64)
```

```

In [221]: # Code from sklearn documentation.
from sklearn.model_selection import learning_curve
from sklearn.model_selection import ShuffleSplit
def plot_learning_curve(estimator, title, X, y, ylim=None, cv=None,
                        n_jobs=1, train_sizes=np.linspace(.1, 1.0, 5)):
    """
    Generate a simple plot of the test and training learning curve.
    """
    plt.figure()
    plt.title(title)
    if ylim is not None:
        plt.ylim(*ylim)
    plt.xlabel("Training examples")
    plt.ylabel("Score")
    train_sizes, train_scores, test_scores = learning_curve(
        estimator, X, y, cv=cv, n_jobs=n_jobs, train_sizes=train_sizes)
    train_scores_mean = np.mean(train_scores, axis=1)
    train_scores_std = np.std(train_scores, axis=1)
    test_scores_mean = np.mean(test_scores, axis=1)
    test_scores_std = np.std(test_scores, axis=1)
    plt.grid()


    plt.fill_between(train_sizes, train_scores_mean - train_scores_std,
                     train_scores_mean + train_scores_std, alpha=0.1,
                     color="r")
    plt.fill_between(train_sizes, test_scores_mean - test_scores_std,
                     test_scores_mean + test_scores_std, alpha=0.1, color="g")
    plt.plot(train_sizes, train_scores_mean, 'o-', color="r",
             label="Training score")
    plt.plot(train_sizes, test_scores_mean, 'o-', color="g",
             label="Cross-validation score")

    plt.legend(loc="best")
    return plt

```



```
In [226]: g = plot_learning_curve(svc.grid.best_estimator_, "SVC Learning Curve", X_train, Y_train, ylim=[0.95, 0.55], cv = 5,
                                train_sizes = [0.05, 0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1])
```



## Logistic Regression

```
In [227]: from sklearn.linear_model import LogisticRegression
```


```
In [228]: lr = Class_Fit(clf = LogisticRegression)
lr.grid_search(parameters = [{'C':np.logspace(-4,6,16)}], Kfold = 17)
lr.grid_fit(X_train, Y_train)
lr.grid_predict(X_test, Y_test)
```

Precision: 85.04 %

```
In [229]: cnf = confusion_matrix(Y_test, lr.predictions)
cnf
```

```
Out[229]: array([[369, 12],
                [ 55, 12]], dtype=int64)
```

```
In [230]: g = plot_learning_curve(lr.grid.best_estimator_, "LogisticRegression Learning Curve", X_train, Y_train, ylim=[.95, 0.75], cv = 5,
                                train_sizes = [0.05, 0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1])
```



## K-Nearest Neighbours :

```
In [231]: from sklearn.neighbors import KNeighborsClassifier
```


```
In [232]: knn = Class_Fit(clf = KNeighborsClassifier)
knn.grid_search(parameters = [{ 'n_neighbors': np.arange(1,50,1) }], Kfold = 10)
knn.grid_fit(X_train, Y_train)
knn.grid_predict(X_test, Y_test)
```

Precision: 83.93 %

```
In [233]: cnf = confusion_matrix(Y_test, knn.predictions)
cnf
```

```
Out[233]: array([[369, 12],
                [ 60,  7]], dtype=int64)
```

```
In [234]: g = plot_learning_curve(knn.grid.best_estimator_, "KNearestNEighbors Learning Curve", X_train, Y_train, ylim=[.9, 0.8],
cv = 5,
                                train_sizes = [0.05, 0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1])
```



## Decision Trees :

```
In [235]: from sklearn.tree import DecisionTreeClassifier
```


```
In [236]: tr = Class_Fit(clf = DecisionTreeClassifier)
tr.grid_search(parameters = [{'criterion':['entropy', 'gini'], 'max_features':['sqrt', 'log2']}], Kfold = 3)
tr.grid_fit(X_train, Y_train)
tr.grid_predict(X_test, Y_test)
```

Precision: 81.03 %

```
In [237]: cnf = confusion_matrix(Y_test, tr.predictions)
cnf
```

```
Out[237]: array([[341, 40],
                [ 45, 22]], dtype=int64)
```

```
In [238]: g = plot_learning_curve(tr.grid.best_estimator_, "DecisionTree Learning Curve", X_train, Y_train, ylim=[1.01, 0.75], cv
= 5,
                                train_sizes = [0.05, 0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1])
```



## Random Forests:

```
In [239]: from sklearn.ensemble import RandomForestClassifier
```

```
In [240]: rf = Class_Fit(clf = RandomForestClassifier)
rf.grid_search(parameters = [{'criterion':['entropy', 'gini'],
                                'max_features':['sqrt', 'log2'], 'n_estimators':[40, 60, 80, 100, 140]}], Kfold = 5)


rf.grid_fit(X_train, Y_train)
rf.grid_predict(X_test, Y_test)
```

Precision: 86.61 %

```
In [258]: cnf = confusion_matrix(Y_test, rf.predictions)
cnf
```

```
Out[258]: array([[385,  5],
                [ 42, 16]], dtype=int64)
```

```
In [241]: g = plot_learning_curve(rf.grid.best_estimator_, "Random Forest Learning Curve", X_train, Y_train, ylim=[1.01, 0.8], cv
= 5,
                                train_sizes = [0.05, 0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1])
```



```
In [242]: from sklearn.ensemble import AdaBoostClassifier
```


```
In [244]: ada = Class_Fit(clf = AdaBoostClassifier)
ada.grid_search(parameters = [{'n_estimators':[20, 30, 40, 50, 60, 70, 80, 90, 100, 120, 130]}], Kfold = 8)
ada.grid_fit(X_train, Y_train)
ada.grid_predict(X_test, Y_test)
```

Precision: 86.16 %

```
In [245]: cnf = confusion_matrix(Y_test, ada.predictions)
cnf
```

```
Out[245]: array([[358, 23],
                [ 39, 28]], dtype=int64)
```

```
In [246]: g = plot_learning_curve(ada.grid.best_estimator_, "AdaBoost Learning Curve", X_train, Y_train, ylim=[1.01, 0.8], cv = 5
,
                                train_sizes = [0.05, 0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1])
```



**Gradient Boosted Decision Trees :**

In [248]: `import xgboost`


In [249]: `gbdt = Class_Fit(clf = xgboost.XGBClassifier)  
gbdt.grid_search(parameters = [{'n_estimators':[20, 30, 40, 50, 60, 70, 80, 90, 100, 120]}], Kfold = 5)  
gbdt.grid_fit(X_train, Y_train)  
gbdt.grid_predict(X_test, Y_test)`

Precision: 85.71 %

In [250]: `cnf = confusion_matrix(Y_test, gbdt.predictions)  
cnf`

Out[250]: `array([[363, 18],  
 [ 46, 21]], dtype=int64)`

In [251]: `g = plot_learning_curve(gbdt.grid.best_estimator_, "GBDT Learning Curve", X_train, Y_train, ylim=[1.01, 0.8], cv = 5,  
 train_sizes = [0.05, 0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1])`



Voting Classifier :

```
In [252]: rf_best = RandomForestClassifier(**rf.grid.best_params_)
gbdt_best = xgboost.XGBClassifier(**gbdt.grid.best_params_)
svc_best = LinearSVC(**svc.grid.best_params_)
tr_best = DecisionTreeClassifier(**tr.grid.best_params_)
knn_best = KNeighborsClassifier(**knn.grid.best_params_)
lr_best = LogisticRegression(**lr.grid.best_params_)
```

```
In [253]: from sklearn.ensemble import VotingClassifier
```

```
In [263]: votingC = VotingClassifier(estimators=[('rf', rf_best), ('gb', gbdt_best), ('knn', knn_best), ('lr', lr_best), ('svc', svc_best), ('tr', tr_best)])
```

```
In [264]: votingC = votingC.fit(X_train, Y_train)
```

```
In [265]: predictions = votingC.predict(X_test)
```

```
In [266]: print("Precision : {:.2f}%".format(100 * accuracy_score(Y_test, predictions)))

Precision : 85.71%
```

```
In [267]: from sklearn.metrics import classification_report
print(classification_report(Y_test, predictions))
```

	precision	recall	f1-score	support
0	0.87	0.98	0.92	381
1	0.59	0.15	0.24	67
accuracy			0.86	448
macro avg	0.73	0.57	0.58	448
weighted avg	0.83	0.86	0.82	448

## Testing Model



In [268]: X

Out[268]:

	renda_mes_media	age	recencia_dias	vinho_montante	frutas_montante	carne_montante	peixe_montante	doces_montante	ouro_montante	p
0	4844	63	58	635	88	546	172	88	88	
1	3862	66	38	11	1	6	2	1	6	
2	5967	55	26	426	49	127	111	21	42	
3	2220	36	26	11	4	20	10	3	5	
4	4857	39	94	173	43	118	46	27	15	
...	...	...	...	...	...	...	...	...	...	
2235	5101	53	46	709	43	182	42	118	247	
2236	5334	74	56	406	0	30	0	0	8	
2237	4748	39	91	908	48	217	32	12	24	
2238	5770	64	8	428	30	214	80	30	61	
2239	4405	66	40	84	3	61	2	1	21	

2240 rows × 22 columns



```
In [269]: # define standard scaler
scaler = StandardScaler()
# transform data
scaled = scaler.fit_transform(X)
```

```
In [270]: predictions = votingC.predict(X)
```

```
In [271]: print("Precision : {:.2f}%".format(100 * accuracy_score(Y, predictions)))
```

Precision : 91.83%